

# Madhur Jaripatke

Roll No. 52

BE A Computer

RMDSSOE, Warje, Pune

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months. Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc.

Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling>

Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix (5 points).

## Importing Libraries

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from keras.models import Sequential
from keras.layers import Dense, Input
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

```

2024-11-06 16:23:11.655811: I tensorflow/core/util/port.cc:153] oneDNN custom
operations are on. You may see slightly different numerical results due to
floating-point round-off errors from different computation orders. To turn t
hem off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.
2024-11-06 16:23:11.656314: I external/local_xla/xla/tsl/cuda/cudart_stub.c
c:32] Could not find cuda drivers on your machine, GPU will not be used.
2024-11-06 16:23:11.659166: I external/local_xla/xla/tsl/cuda/cudart_stub.c
c:32] Could not find cuda drivers on your machine, GPU will not be used.
2024-11-06 16:23:11.666761: E external/local_xla/xla/stream_executor/cuda/cu
da_fft.cc:485] Unable to register cuFFT factory: Attempting to register fac
tory for plugin cuFFT when one has already been registered
2024-11-06 16:23:11.678891: E external/local_xla/xla/stream_executor/cuda/cu
da_dnn.cc:8454] Unable to register cuDNN factory: Attempting to register fac
tory for plugin cuDNN when one has already been registered
2024-11-06 16:23:11.682973: E external/local_xla/xla/stream_executor/cuda/cu
da_blas.cc:1452] Unable to register cuBLAS factory: Attempting to register f
actory for plugin cuBLAS when one has already been registered
2024-11-06 16:23:11.692873: I tensorflow/core/platform/cpu_feature_guard.cc:
210] This TensorFlow binary is optimized to use available CPU instructions i
n performance-critical operations.
To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other
operations, rebuild TensorFlow with the appropriate compiler flags.
2024-11-06 16:23:12.418863: W tensorflow/compiler/tf2tensorrt/utils/py_util
s.cc:38] TF-TRT Warning: Could not find TensorRT

```

## Loading the Dataset

```
In [2]: df = pd.read_csv('./Datasets/churn_modelling.csv')
df.head()
```

```
Out[2]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

```
In [3]: df.shape
```

```
Out[3]: (10000, 14)
```

```
In [4]: df.describe
```

```

Out[4]: <bound method NDFrame.describe of
reditScore Geography Gender Age \
0 1 15634602 Hargrave 619 France Female 42
1 2 15647311 Hill 608 Spain Female 41
2 3 15619304 Onio 502 France Female 42
3 4 15701354 Boni 699 France Female 39
4 5 15737888 Mitchell 850 Spain Female 43
...
9995 9996 15606229 Obijiaku 771 France Male 39
9996 9997 15569892 Johnstone 516 France Male 35
9997 9998 15584532 Liu 709 France Female 36
9998 9999 15682355 Sabbatini 772 Germany Male 42
9999 10000 15628319 Walker 792 France Female 28

Tenure Balance NumOfProducts HasCrCard IsActiveMember \
0 2 0.00 1 1 1
1 1 83807.86 1 0 1
2 8 159660.80 3 1 0
3 1 0.00 2 0 0
4 2 125510.82 1 1 1
...
9995 5 0.00 2 1 0
9996 10 57369.61 1 1 1
9997 7 0.00 1 0 1
9998 3 75075.31 2 1 0
9999 4 130142.79 1 1 0

EstimatedSalary Exited
0 101348.88 1
1 112542.58 0
2 113931.57 1
3 93826.63 0
4 79084.10 0
...
9995 96270.64 0
9996 101699.77 0
9997 42085.58 1
9998 92888.52 1
9999 38190.78 0

```

```
[10000 rows x 14 columns]>
```

## Exploratory Data Analysis (EDA)

```

In [5]: df.isnull()
df.isnull().sum()

```

```
Out[5]: RowNumber      0
        CustomerId     0
        Surname        0
        CreditScore     0
        Geography      0
        Gender         0
        Age            0
        Tenure         0
        Balance        0
        NumOfProducts  0
        HasCrCard      0
        IsActiveMember 0
        EstimatedSalary 0
        Exited         0
        dtype: int64
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender               10000 non-null  object
6   Age                  10000 non-null  int64
7   Tenure               10000 non-null  int64
8   Balance              10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard            10000 non-null  int64
11  IsActiveMember       10000 non-null  int64
12  EstimatedSalary      10000 non-null  float64
13  Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [7]: df.dtypes
```

```
Out[7]: RowNumber      int64
        CustomerId     int64
        Surname        object
        CreditScore     int64
        Geography      object
        Gender          object
        Age             int64
        Tenure          int64
        Balance         float64
        NumOfProducts  int64
        HasCrCard       int64
        IsActiveMember  int64
        EstimatedSalary float64
        Exited          int64
        dtype: object
```

```
In [8]: df.columns
```

```
Out[8]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
              'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
              'IsActiveMember', 'EstimatedSalary', 'Exited'],
              dtype='object')
```

## Data Preprocessing

```
In [9]: df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis = 1)
        df.head()
```

```
Out[9]:
```


	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
0	619	France	Female	42	2	0.00	1
1	608	Spain	Female	41	1	83807.86	1
2	502	France	Female	42	8	159660.80	3
3	699	France	Female	39	1	0.00	2
4	850	Spain	Female	43	2	125510.82	1

## Visualisation

```
In [10]: def visualization(x, y, xlabel):
        plt.figure(figsize=(10,5))
        plt.hist([x, y], color=['yellow', 'blue'], label = ['exited', 'not_exited'])
        plt.xlabel(xlabel, fontsize=20)
        plt.ylabel('No. of customers', fontsize=20)
        plt.legend()
```


```
In [11]: df_churn_exited = df[df['Exited']==1]['Tenure']
        df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
In [12]: visualization(df_churn_exited, df_churn_not_exited, 'Tenure')
```



```
In [13]: df_churn_exited2 = df[df['Exited']==1]['Age']
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
In [14]: visualization(df_churn_exited2, df_churn_not_exited2, 'Age')
```



```
In [15]: x = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
                'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
```

```
In [16]: df = pd.concat([df, gender, states], axis = 1)
df.head()
```

```
Out[16]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
0	619	France	Female	42	2	0.00	1
1	608	Spain	Female	41	1	83807.86	1
2	502	France	Female	42	8	159660.80	3
3	699	France	Female	39	1	0.00	2
4	850	Spain	Female	43	2	125510.82	1

```
In [17]: x = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'IsActive']]
y = df['Exited']
```

## Splitting the Dataset

```
In [18]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3)
```

## Normalising the Data

```
In [19]: sc = StandardScaler()
```

```
In [20]: x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
In [21]: print('Training Data:\n',x_train)
```

```
Training Data:
[[ 0.96818741 -0.56596928 -1.39661597 ... -1.09823226 -0.57603061
  -0.56789208]
 [ 1.19521436 -0.18753916  0.68426206 ...  0.91055421 -0.57603061
  1.76089794]
 [-0.59004305 -0.37675422 -1.04980297 ...  0.91055421 -0.57603061
  -0.56789208]
 ...
 [ 0.27678714  1.98843404  1.72470107 ... -1.09823226 -0.57603061
  -0.56789208]
 [-0.38365491  1.13696627 -0.70298996 ...  0.91055421 -0.57603061
  -0.56789208]
 [-0.70355653  0.47471356  1.03107506 ... -1.09823226 -0.57603061
  -0.56789208]]
```

```
In [22]: print('Testing Data:\n',x_test)
```

```

Testing Data:
[[-2.04507944e+00 -9.29316256e-02  1.03107506e+00 ... -1.09823226e+00
  1.73601886e+00 -5.67892082e-01]
[-1.15350332e-01 -1.79586717e+00 -3.56176956e-01 ... -1.09823226e+00
 -5.76030610e-01 -5.67892082e-01]
[-1.66326138e+00  1.67590482e-03 -1.39661597e+00 ...  9.10554206e-01
 -5.76030610e-01 -5.67892082e-01]
...
[ 1.22617258e+00 -1.41743705e+00  1.37788807e+00 ...  9.10554206e-01
 -5.76030610e-01 -5.67892082e-01]
[ 3.94407736e-02 -9.29316256e-02  1.37788807e+00 ...  9.10554206e-01
 -5.76030610e-01  1.76089794e+00]
[ 1.25713080e+00  1.23157380e+00 -1.04980297e+00 ...  9.10554206e-01
 -5.76030610e-01  1.76089794e+00]]

```

## Building the Neural Network Model

```
In [23]: classifier = Sequential()
```

```
In [24]: classifier.add(Input(shape=(10,)))
```

```
In [25]: classifier.add(Dense(units = 6, kernel_initializer = 'he_uniform', activation = 'relu'))
classifier.add(Dense(units = 6, kernel_initializer = 'he_uniform', activation = 'relu'))
classifier.add(Dense(units = 1, kernel_initializer = 'glorot_uniform', activation = 'sigmoid'))
classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics = ['accuracy'])
```

```
In [26]: classifier.summary()
```

**Model: "sequential"**

Layer (type)	Output Shape	Param #
dense (Dense)	(None, 6)	61
dense_1 (Dense)	(None, 6)	42
dense_2 (Dense)	(None, 1)	7

**Total params: 115** (460.00 B)


**Trainable params: 115** (460.00 B)


**Non-trainable params: 0** (0.00 B)


## Training the Model


```
In [27]: classifier.fit(x_train, y_train, batch_size = 10, epochs = 50)
```





Epoch 1/50  
**700/700**  1s 1ms/step - accuracy: 0.8044 - loss: 0.4926


Epoch 2/50  
**700/700**  1s 1ms/step - accuracy: 0.8125 - loss: 0.4353


Epoch 3/50  
**700/700**  1s 984us/step - accuracy: 0.8244 - loss: 0.4096


Epoch 4/50  
**700/700**  1s 880us/step - accuracy: 0.8353 - loss: 0.3922


Epoch 5/50  
**700/700**  1s 845us/step - accuracy: 0.8432 - loss: 0.3796


Epoch 6/50  
**700/700**  1s 856us/step - accuracy: 0.8422 - loss: 0.3821


Epoch 7/50  
**700/700**  1s 877us/step - accuracy: 0.8514 - loss: 0.3604


Epoch 8/50  
**700/700**  1s 844us/step - accuracy: 0.8516 - loss: 0.3587


Epoch 9/50  
**700/700**  1s 1ms/step - accuracy: 0.8466 - loss: 0.3601


Epoch 10/50  
**700/700**  1s 903us/step - accuracy: 0.8532 - loss: 0.3516


Epoch 11/50  
**700/700**  1s 1ms/step - accuracy: 0.8553 - loss: 0.3512


Epoch 12/50  
**700/700**  1s 1ms/step - accuracy: 0.8531 - loss: 0.3550


Epoch 13/50  
**700/700**  1s 2ms/step - accuracy: 0.8571 - loss: 0.3474


Epoch 14/50  
**700/700**  1s 941us/step - accuracy: 0.8578 - loss: 0.3482


Epoch 15/50  
**700/700**  1s 876us/step - accuracy: 0.8645 - loss: 0.3376


Epoch 16/50  
**700/700**  1s 946us/step - accuracy: 0.8579 - loss: 0.3509


Epoch 17/50  
**700/700**  1s 926us/step - accuracy: 0.8607 - loss: 0.3357


Epoch 18/50  
**700/700**  1s 937us/step - accuracy: 0.8558 - loss: 0.3403


Epoch 19/50  
**700/700**  1s 905us/step - accuracy: 0.8570 - loss: 0.3469


Epoch 20/50  
**700/700**  1s 930us/step - accuracy: 0.8551 - loss: 0.3494


Epoch 21/50  
**700/700**  1s 895us/step - accuracy: 0.8600 - loss: 0.3393


Epoch 22/50  
**700/700**  1s 958us/step - accuracy: 0.8552 - loss: 0.3418


Epoch 23/50  
**700/700**  1s 830us/step - accuracy: 0.8575 - loss: 0.3437

Epoch 24/50  
**700/700**  1s 836us/step - accuracy: 0.8582 - loss: 0.3385

Epoch 25/50  
**700/700**  1s 836us/step - accuracy: 0.8626 - loss: 0.3301

Epoch 26/50  
**700/700**  1s 1ms/step - accuracy: 0.8586 - loss: 0.3398

Epoch 27/50  
**700/700**  1s 1ms/step - accuracy: 0.8522 - loss: 0.3455

Epoch 28/50  
**700/700**  1s 1ms/step - accuracy: 0.8607 - loss: 0.3337

```

Epoch 29/50
700/700 ————— 1s 1ms/step - accuracy: 0.8582 - loss: 0.3372
Epoch 30/50
700/700 ————— 1s 1ms/step - accuracy: 0.8584 - loss: 0.3320
Epoch 31/50
700/700 ————— 1s 1ms/step - accuracy: 0.8602 - loss: 0.3390
Epoch 32/50
700/700 ————— 1s 1ms/step - accuracy: 0.8622 - loss: 0.3389
Epoch 33/50
700/700 ————— 1s 1ms/step - accuracy: 0.8631 - loss: 0.3287
Epoch 34/50
700/700 ————— 1s 1ms/step - accuracy: 0.8557 - loss: 0.3428
Epoch 35/50
700/700 ————— 1s 2ms/step - accuracy: 0.8612 - loss: 0.3325
Epoch 36/50
700/700 ————— 1s 2ms/step - accuracy: 0.8590 - loss: 0.3315
Epoch 37/50
700/700 ————— 1s 2ms/step - accuracy: 0.8608 - loss: 0.3407
Epoch 38/50
700/700 ————— 1s 1ms/step - accuracy: 0.8584 - loss: 0.3339
Epoch 39/50
700/700 ————— 1s 1ms/step - accuracy: 0.8550 - loss: 0.3505
Epoch 40/50
700/700 ————— 1s 1ms/step - accuracy: 0.8572 - loss: 0.3445
Epoch 41/50
700/700 ————— 1s 2ms/step - accuracy: 0.8500 - loss: 0.3516
Epoch 42/50
700/700 ————— 2s 2ms/step - accuracy: 0.8657 - loss: 0.3259
Epoch 43/50
700/700 ————— 1s 2ms/step - accuracy: 0.8583 - loss: 0.3440
Epoch 44/50
700/700 ————— 1s 2ms/step - accuracy: 0.8632 - loss: 0.3321
Epoch 45/50
700/700 ————— 1s 2ms/step - accuracy: 0.8620 - loss: 0.3347
Epoch 46/50
700/700 ————— 1s 1ms/step - accuracy: 0.8565 - loss: 0.3462
Epoch 47/50
700/700 ————— 1s 1ms/step - accuracy: 0.8512 - loss: 0.3470
Epoch 48/50
700/700 ————— 1s 1ms/step - accuracy: 0.8602 - loss: 0.3408
Epoch 49/50
700/700 ————— 2s 2ms/step - accuracy: 0.8549 - loss: 0.3435
Epoch 50/50
700/700 ————— 1s 1ms/step - accuracy: 0.8623 - loss: 0.3282

```

Out[27]: <keras.src.callbacks.history.History at 0x7f07d9374470>

## Evaluating the Model

```

In [28]: y_pred = classifier.predict(x_test)
         y_pred = (y_pred > 0.5)

```

```

94/94 ————— 0s 1ms/step

```

```
In [29]: cm = confusion_matrix(y_test, y_pred)
print('Confusion Matrix:\n', cm)
```


```
Confusion Matrix:
[[2285  108]
 [ 307  300]]
```

```
In [30]: acc = accuracy_score(y_test, y_pred)
print('Accuracy Score:', acc)
```

```
Accuracy Score: 0.8616666666666667
```

```
In [31]: plt.figure(figsize=(10,7))
sns.heatmap(cm, annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

```
Out[31]: Text(95.7222222222221, 0.5, 'Truth')
```



```
In [32]: print('Classification Report:\n', classification_report(y_test, y_pred))
```

Classification Report:				
	precision	recall	f1-score	support
0	0.88	0.95	0.92	2393
1	0.74	0.49	0.59	607
accuracy			0.86	3000
macro avg	0.81	0.72	0.75	3000
weighted avg	0.85	0.86	0.85	3000