

Madhur Jaripatke

Roll No. 52

BE A Computer

RMDSSOE, Warje, Pune

Given a bank customer, build a neural network-based classifier that can determine whether they will leave or not in the next 6 months. Dataset Description: The case study is from an open-source dataset from Kaggle. The dataset contains 10,000 sample points with 14 distinct features such as CustomerId, CreditScore, Geography, Gender, Age, Tenure, Balance, etc. Link to the Kaggle project: <https://www.kaggle.com/barelydedicated/bank-customer-churn-modeling>

Perform following steps:

1. Read the dataset.
2. Distinguish the feature and target set and divide the data set into training and test sets.
3. Normalize the train and test data.
4. Initialize and build the model. Identify the points of improvement and implement the same.
5. Print the accuracy score and confusion matrix (5 points).

```
In [1]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from keras.models import Sequential
from keras.layers import Dense
from sklearn.metrics import confusion_matrix, accuracy_score, classification_report
```

2024-10-24 22:04:27.219031: I tensorflow/core/util/port.cc:153] oneDNN custom operations are on. You may see slightly different numerical results due to floating-point round-off errors from different computation orders. To turn them off, set the environment variable `TF_ENABLE_ONEDNN_OPTS=0`.

2024-10-24 22:04:27.219821: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32] Could not find cuda drivers on your machine, GPU will not be used.

2024-10-24 22:04:27.223789: I external/local_xla/xla/tsl/cuda/cudart_stub.cc:32] Could not find cuda drivers on your machine, GPU will not be used.

2024-10-24 22:04:27.233450: E external/local_xla/xla/stream_executor/cuda/cuda_fft.cc:485] Unable to register cuFFT factory: Attempting to register factory for plugin cuFFT when one has already been registered

2024-10-24 22:04:27.247799: E external/local_xla/xla/stream_executor/cuda/cuda_dnn.cc:8454] Unable to register cuDNN factory: Attempting to register factory for plugin cuDNN when one has already been registered

2024-10-24 22:04:27.253778: E external/local_xla/xla/stream_executor/cuda/cuda_blas.cc:1452] Unable to register cuBLAS factory: Attempting to register factory for plugin cuBLAS when one has already been registered

2024-10-24 22:04:27.267679: I tensorflow/core/platform/cpu_feature_guard.cc:210] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.

To enable the following instructions: AVX2 AVX512F AVX512_VNNI FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

2024-10-24 22:04:28.514490: W tensorflow/compiler/tf2tensorrt/utils/py_utils.cc:38] TF-TRT Warning: Could not find TensorRT

```
In [2]: df = pd.read_csv('./Datasets/churn_modelling.csv')
df.head()
```

```
Out[2]:
```

	RowNumber	CustomerId	Surname	CreditScore	Geography	Gender	Age
0	1	15634602	Hargrave	619	France	Female	42
1	2	15647311	Hill	608	Spain	Female	41
2	3	15619304	Onio	502	France	Female	42
3	4	15701354	Boni	699	France	Female	39
4	5	15737888	Mitchell	850	Spain	Female	43

```
In [3]: df.shape
```

```
Out[3]: (10000, 14)
```

```
In [4]: df.describe
```

```
Out[4]: <bound method NDFrame.describe of
reditScore Geography Gender Age \
0 1 15634602 Hargrave 619 France Female 42
1 2 15647311 Hill 608 Spain Female 41
2 3 15619304 Onio 502 France Female 42
3 4 15701354 Boni 699 France Female 39
4 5 15737888 Mitchell 850 Spain Female 43
...
9995 9996 15606229 Obijiaku 771 France Male 39
9996 9997 15569892 Johnstone 516 France Male 35
9997 9998 15584532 Liu 709 France Female 36
9998 9999 15682355 Sabbatini 772 Germany Male 42
9999 10000 15628319 Walker 792 France Female 28
```

```
Tenure Balance NumOfProducts HasCrCard IsActiveMember \
0 2 0.00 1 1 1
1 1 83807.86 1 0 1
2 8 159660.80 3 1 0
3 1 0.00 2 0 0
4 2 125510.82 1 1 1
...
9995 5 0.00 2 1 0
9996 10 57369.61 1 1 1
9997 7 0.00 1 0 1
9998 3 75075.31 2 1 0
9999 4 130142.79 1 1 0
```

```
EstimatedSalary Exited
0 101348.88 1
1 112542.58 0
2 113931.57 1
3 93826.63 0
4 79084.10 0
...
9995 96270.64 0
9996 101699.77 0
9997 42085.58 1
9998 92888.52 1
9999 38190.78 0
```

```
[10000 rows x 14 columns]>
```

```
In [5]: df.isnull()
df.isnull().sum()
```

```
Out[5]: RowNumber      0
        CustomerId     0
        Surname        0
        CreditScore     0
        Geography      0
        Gender         0
        Age            0
        Tenure         0
        Balance        0
        NumOfProducts  0
        HasCrCard      0
        IsActiveMember 0
        EstimatedSalary 0
        Exited         0
        dtype: int64
```

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 10000 entries, 0 to 9999
Data columns (total 14 columns):
#   Column                Non-Null Count  Dtype
---  ---
0   RowNumber             10000 non-null  int64
1   CustomerId            10000 non-null  int64
2   Surname               10000 non-null  object
3   CreditScore           10000 non-null  int64
4   Geography             10000 non-null  object
5   Gender               10000 non-null  object
6   Age                  10000 non-null  int64
7   Tenure               10000 non-null  int64
8   Balance              10000 non-null  float64
9   NumOfProducts        10000 non-null  int64
10  HasCrCard            10000 non-null  int64
11  IsActiveMember       10000 non-null  int64
12  EstimatedSalary      10000 non-null  float64
13  Exited               10000 non-null  int64
dtypes: float64(2), int64(9), object(3)
memory usage: 1.1+ MB
```

```
In [7]: df.dtypes
```

```
Out[7]: RowNumber      int64
CustomerId    int64
Surname       object
CreditScore   int64
Geography     object
Gender        object
Age           int64
Tenure        int64
Balance       float64
NumOfProducts int64
HasCrCard     int64
IsActiveMember int64
EstimatedSalary float64
Exited        int64
dtype: object
```

```
In [8]: df.columns
```

```
Out[8]: Index(['RowNumber', 'CustomerId', 'Surname', 'CreditScore', 'Geography',
              'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'HasCrCard',
              'IsActiveMember', 'EstimatedSalary', 'Exited'],
              dtype='object')
```

```
In [9]: df = df.drop(['RowNumber', 'Surname', 'CustomerId'], axis = 1)
df.head()
```


```
Out[9]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
0	619	France	Female	42	2	0.00	1
1	608	Spain	Female	41	1	83807.86	1
2	502	France	Female	42	8	159660.80	3
3	699	France	Female	39	1	0.00	2
4	850	Spain	Female	43	2	125510.82	1

```
In [10]: def visualization(x, y, xlabel):
plt.figure(figsize=(10,5))
plt.hist([x, y], color=['yellow', 'blue'], label = ['exited', 'not_exited'])
plt.xlabel(xlabel, fontsize=20)
plt.ylabel('No. of customers', fontsize=20)
plt.legend()
```


```
In [11]: df_churn_exited = df[df['Exited']==1]['Tenure']
df_churn_not_exited = df[df['Exited']==0]['Tenure']
```

```
In [12]: visualization(df_churn_exited, df_churn_not_exited, 'Tenure')
```



```
In [13]: df_churn_exited2 = df[df['Exited']==1]['Age']
df_churn_not_exited2 = df[df['Exited']==0]['Age']
```

```
In [14]: visualization(df_churn_exited2, df_churn_not_exited2, 'Age')
```



```
In [15]: x = df[['CreditScore', 'Gender', 'Age', 'Tenure', 'Balance', 'NumOfProducts',
                'HasCrCard', 'IsActiveMember', 'EstimatedSalary']]
states = pd.get_dummies(df['Geography'], drop_first = True)
gender = pd.get_dummies(df['Gender'], drop_first = True)
```

```
In [16]: df = pd.concat([df, gender, states], axis = 1)
df.head()
```

```
Out[16]:
```

	CreditScore	Geography	Gender	Age	Tenure	Balance	NumOfProducts
0	619	France	Female	42	2	0.00	1
1	608	Spain	Female	41	1	83807.86	1
2	502	France	Female	42	8	159660.80	3
3	699	France	Female	39	1	0.00	2
4	850	Spain	Female	43	2	125510.82	1

```
In [17]: x = df[['CreditScore', 'Age', 'Tenure', 'Balance', 'NumOfProducts', 'IsActive']]
y = df['Exited']
```

```
In [18]: x_train, x_test, y_train, y_test = train_test_split(x, y, test_size = 0.3)
```

```
In [19]: sc = StandardScaler()
```

```
In [20]: x_train = sc.fit_transform(x_train)
x_test = sc.transform(x_test)
```

```
In [21]: print('Training Data:\n',x_train)
```

```
Training Data:
[[-0.25855342 -0.18536388 -0.33750179 ... -1.10522259  1.72416885
 -0.56811207]
 [ 1.06208156  0.00809775  1.04894922 ... -1.10522259  1.72416885
 -0.56811207]
 [ 0.31474549  0.20155938 -1.72395279 ... -1.10522259 -0.5799896
 -0.56811207]
 ...
 [ 0.73448218  0.10482857  0.35572371 ... -1.10522259  1.72416885
 -0.56811207]
 [ 1.2463562   1.16886755  1.39556197 ... -1.10522259  1.72416885
 -0.56811207]
 [-2.10129989 -1.34613368  0.35572371 ...  0.90479511 -0.5799896
 -0.56811207]]
```

```
In [22]: print('Testing Data:\n',x_test)
```

```
Testing Data:
[[-0.35069074 -0.18536388  1.39556197 ...  0.90479511 -0.5799896
 -0.56811207]
 [-1.48705107 -0.37882552  1.04894922 ...  0.90479511 -0.5799896
  1.76021608]
 [ 2.04487968 -0.2820947   0.70233647 ...  0.90479511  1.72416885
 -0.56811207]
 ...
 [ 0.10999588 -0.47555633 -1.72395279 ... -1.10522259 -0.5799896
 -0.56811207]
 [ 0.30450801 -0.18536388  1.74217472 ... -1.10522259 -0.5799896
 -0.56811207]
 [-0.95470209  0.2982902   0.35572371 ...  0.90479511 -0.5799896
 -0.56811207]]
```

```
In [23]: classifier = Sequential()
```

```
In [24]: classifier.add(Dense(units = 6, kernel_initializer = 'he_uniform', activation='sigmoid'))
classifier.add(Dense(units = 6, kernel_initializer = 'he_uniform', activation='sigmoid'))
classifier.add(Dense(units = 1, kernel_initializer = 'glorot_uniform', activation='sigmoid'))
classifier.compile(optimizer = 'adam', loss = 'binary_crossentropy', metrics=['accuracy'])
```

/home/madhurj20/.local/lib/python3.12/site-packages/keras/src/layers/core/dense.py:87: UserWarning: Do not pass an `input_shape`/`input_dim` argument to a layer. When using Sequential models, prefer using an `Input(shape)` object as the first layer in the model instead.

```
super().__init__(activity_regularizer=activity_regularizer, **kwargs)
```

```
In [25]: classifier.summary()
```

Model: "sequential"





























Layer (type)	Output Shape	Param #
dense (Dense)	(None, 6)	6
dense_1 (Dense)	(None, 6)	6
dense_2 (Dense)	(None, 1)	1

Total params: 115 (460.00 B)

Trainable params: 115 (460.00 B)

Non-trainable params: 0 (0.00 B)

```
In [26]: classifier.fit(x_train, y_train, batch_size = 10, epochs = 50)
```


Epoch 1/50			
700/700		2s 1ms/step	- accuracy: 0.7247 - loss: 0.5849
Epoch 2/50			
700/700		1s 1ms/step	- accuracy: 0.8061 - loss: 0.4421
Epoch 3/50			
700/700		1s 1ms/step	- accuracy: 0.8136 - loss: 0.4182
Epoch 4/50			
700/700		1s 1ms/step	- accuracy: 0.8304 - loss: 0.3938
Epoch 5/50			
700/700		1s 1ms/step	- accuracy: 0.8349 - loss: 0.3846
Epoch 6/50			
700/700		1s 1ms/step	- accuracy: 0.8385 - loss: 0.3726
Epoch 7/50			
700/700		1s 1ms/step	- accuracy: 0.8504 - loss: 0.3520
Epoch 8/50			
700/700		1s 1ms/step	- accuracy: 0.8521 - loss: 0.3589
Epoch 9/50			
700/700		1s 1ms/step	- accuracy: 0.8591 - loss: 0.3498
Epoch 10/50			
700/700		1s 1ms/step	- accuracy: 0.8552 - loss: 0.3442
Epoch 11/50			
700/700		1s 1ms/step	- accuracy: 0.8570 - loss: 0.3499
Epoch 12/50			
700/700		1s 1ms/step	- accuracy: 0.8606 - loss: 0.3435
Epoch 13/50			
700/700		1s 1ms/step	- accuracy: 0.8533 - loss: 0.3508
Epoch 14/50			
700/700		1s 1ms/step	- accuracy: 0.8572 - loss: 0.3427
Epoch 15/50			
700/700		1s 1ms/step	- accuracy: 0.8560 - loss: 0.3404
Epoch 16/50			
700/700		1s 1ms/step	- accuracy: 0.8582 - loss: 0.3394
Epoch 17/50			
700/700		1s 967us/step	- accuracy: 0.8659 - loss: 0.3279
Epoch 18/50			
700/700		1s 996us/step	- accuracy: 0.8664 - loss: 0.3179
Epoch 19/50			
700/700		1s 1ms/step	- accuracy: 0.8630 - loss: 0.3274
Epoch 20/50			
700/700		1s 1ms/step	- accuracy: 0.8603 - loss: 0.3418
Epoch 21/50			
700/700		1s 1ms/step	- accuracy: 0.8630 - loss: 0.3286
Epoch 22/50			
700/700		1s 1ms/step	- accuracy: 0.8563 - loss: 0.3371
Epoch 23/50			
700/700		1s 1ms/step	- accuracy: 0.8549 - loss: 0.3363
Epoch 24/50			
700/700		1s 1ms/step	- accuracy: 0.8547 - loss: 0.3454
Epoch 25/50			
700/700		1s 1ms/step	- accuracy: 0.8482 - loss: 0.3524
Epoch 26/50			
700/700		1s 1ms/step	- accuracy: 0.8600 - loss: 0.3393
Epoch 27/50			
700/700		1s 1ms/step	- accuracy: 0.8607 - loss: 0.3342
Epoch 28/50			
700/700		1s 1ms/step	- accuracy: 0.8637 - loss: 0.3305

```

Epoch 29/50
700/700 ————— 1s 1ms/step - accuracy: 0.8624 - loss: 0.3296
Epoch 30/50
700/700 ————— 1s 1ms/step - accuracy: 0.8633 - loss: 0.3279
Epoch 31/50
700/700 ————— 1s 1ms/step - accuracy: 0.8646 - loss: 0.3298
Epoch 32/50
700/700 ————— 1s 1ms/step - accuracy: 0.8602 - loss: 0.3331
Epoch 33/50
700/700 ————— 1s 1ms/step - accuracy: 0.8624 - loss: 0.3312
Epoch 34/50
700/700 ————— 1s 1ms/step - accuracy: 0.8614 - loss: 0.3294
Epoch 35/50
700/700 ————— 1s 1ms/step - accuracy: 0.8526 - loss: 0.3440
Epoch 36/50
700/700 ————— 1s 1ms/step - accuracy: 0.8642 - loss: 0.3304
Epoch 37/50
700/700 ————— 1s 1ms/step - accuracy: 0.8522 - loss: 0.3449
Epoch 38/50
700/700 ————— 1s 1ms/step - accuracy: 0.8548 - loss: 0.3444
Epoch 39/50
700/700 ————— 1s 1ms/step - accuracy: 0.8596 - loss: 0.3384
Epoch 40/50
700/700 ————— 1s 1ms/step - accuracy: 0.8659 - loss: 0.3312
Epoch 41/50
700/700 ————— 1s 1ms/step - accuracy: 0.8650 - loss: 0.3241
Epoch 42/50
700/700 ————— 1s 1ms/step - accuracy: 0.8564 - loss: 0.3475
Epoch 43/50
700/700 ————— 1s 1ms/step - accuracy: 0.8604 - loss: 0.3286
Epoch 44/50
700/700 ————— 1s 1ms/step - accuracy: 0.8682 - loss: 0.3208
Epoch 45/50
700/700 ————— 1s 1ms/step - accuracy: 0.8537 - loss: 0.3402
Epoch 46/50
700/700 ————— 1s 1ms/step - accuracy: 0.8629 - loss: 0.3341
Epoch 47/50
700/700 ————— 1s 1ms/step - accuracy: 0.8620 - loss: 0.3321
Epoch 48/50
700/700 ————— 1s 1ms/step - accuracy: 0.8540 - loss: 0.3443
Epoch 49/50
700/700 ————— 1s 1ms/step - accuracy: 0.8569 - loss: 0.3421
Epoch 50/50
700/700 ————— 1s 1ms/step - accuracy: 0.8697 - loss: 0.3246

```

Out[26]: <keras.src.callbacks.history.History at 0x7f248e49b500>

```
In [27]: y_pred = classifier.predict(x_test)
         y_pred = (y_pred > 0.5)
```

```
94/94 ————— 0s 1ms/step
```

```
In [28]: cm = confusion_matrix(y_test, y_pred)
         print('Confusion Matrix:\n', cm)
```

Confusion Matrix:


```
[[2280  97]
 [ 320 303]]
```

```
In [29]: acc = accuracy_score(y_test, y_pred)
print('Accuracy Score:', acc)
```

Accuracy Score: 0.861

```
In [30]: plt.figure(figsize=(10,7))
sns.heatmap(cm, annot = True)
plt.xlabel('Predicted')
plt.ylabel('Truth')
```

Out[30]: Text(95.7222222222221, 0.5, 'Truth')



```
In [31]: print('Classification Report:\n',classification_report(y_test, y_pred))
```

Classification Report:

	precision	recall	f1-score	support
0	0.88	0.96	0.92	2377
1	0.76	0.49	0.59	623
accuracy			0.86	3000
macro avg	0.82	0.72	0.75	3000
weighted avg	0.85	0.86	0.85	3000